

November 2024

R Shiny Masterclass Series - Advanced

Session 5

Advanced data sources and performance of R Shiny apps



EPI-interactive

Agenda

Session 1 | October 30th | Introduction, recap and responsive interfaces in R Shiny

Session 2 | November 1st | Advanced reactivity and UX considerations

Session 3 | November 5th | Useful R packages to extend core Shiny functionality

Session 4 | November 6th | Managing complexity: modularizing with the module pattern

Session 5 | November 8th | Advanced data sources and processing

Session 6 | November 12th | Automated report generation

Session 7 | November 13th | User authentication, Extended exercise

Session 8 | November 15th | AI Tools, Programming sins and how to avoid them

Today

- Namespace Tips
- Databases in Shiny
- Performance considerations in Shiny apps

Debugging Namespaces

```
download_ui <- function(id){  
  ns <- NS(id)  
  message(ns('my_id'))  
}
```

```
download_server <- function(id, data) {  
  moduleServer(  
    id,  
    function(input, output, session) {  
      message(session$ns('my_id'))  
    }  
  )  
}
```

```
> runApp('challenge')
```

```
Listening on http://127.0.0.1:4713  
table-download_data-my_id  
region-region_summary_table-download_data-my_id  
table-table-download_data-my_id  
region-region_summary_table-region-region_summary_table-download_data-my_id
```

```
> runApp('challenge')
```

```
Listening on http://127.0.0.1:4713  
table-download_data-my_id  
region-region_summary_table-download_data-my_id  
table-download_data-my_id  
region-region_summary_table-download_data-my_id
```

Databases in Shiny

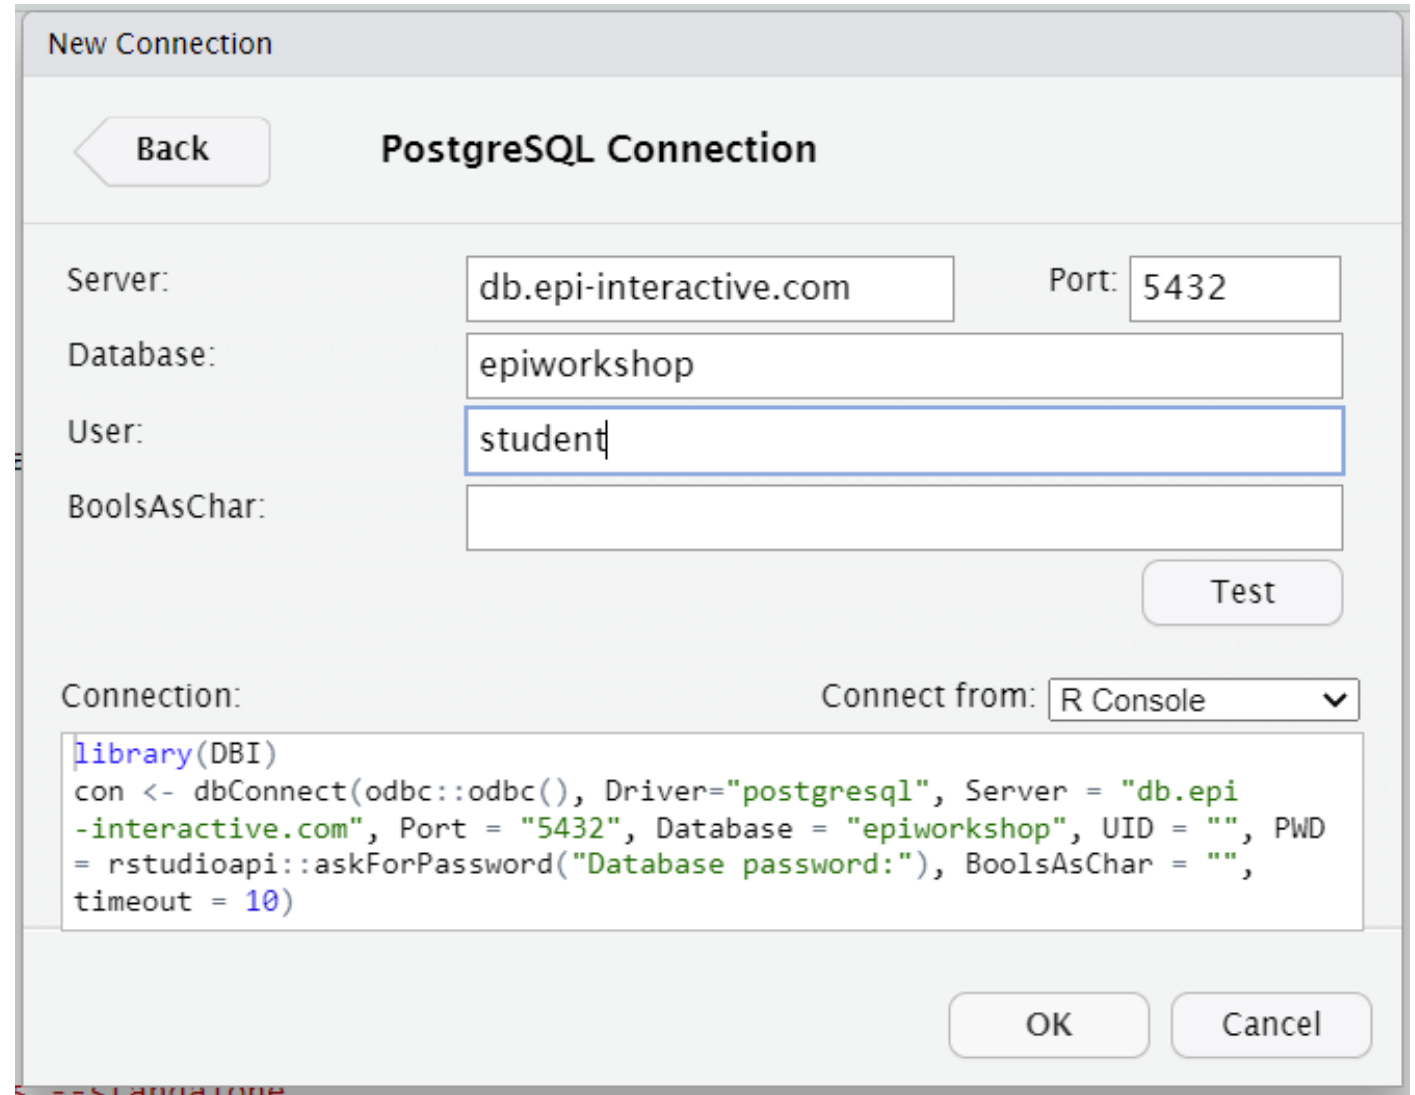
When to use databases over files

- Well defined, structured data
- Need to frequently insert/update
 - Safety against multiple users doing so at the same time
- More complexity to the data
- Remote data

Integrating to our app: R Libraries

- [odbc](#): Allows us to use the correct database drivers for our connection
- [DBI](#): Connection management for singular database connections
- [pool](#): Connection management for the database – multiple connections

Database connections in Posit Cloud



The screenshot shows a 'New Connection' dialog box with a 'PostgreSQL Connection' tab. The dialog includes fields for Server, Database, User, and BoolAsChar, along with a Port field. A 'Test' button is located to the right of the BoolAsChar field. Below these fields, there is a 'Connection:' label and a 'Connect from:' dropdown menu set to 'R Console'. A text area contains R code for connecting to the database. At the bottom, there are 'OK' and 'Cancel' buttons.

New Connection

Back PostgreSQL Connection

Server: db.epi-interactive.com Port: 5432

Database: epiworkshop

User: student

BoolAsChar:

Test

Connection: Connect from: R Console

```
library(DBI)
con <- dbConnect(odbc::odbc(), Driver="postgresql", Server = "db.epi-interactive.com", Port = "5432", Database = "epiworkshop", UID = "", PWD = rstudioapi::askForPassword("Database password:"), BoolAsChar = "", timeout = 10)
```

OK Cancel

Integrating to our app: global.R

```
library(odbc)
library(DBI)

dbConn <- dbConnect(
  RPostgres::Postgres(),
  host = "db.epi-interactive.com",
  port = 5432,
  dbname = "epiworkshop",
  user = "student",
  password = "student"
)
```

Where to create the connection?

- Global environment
 - One connection can be shared by multiple user sessions
 - Need to leave an open connection – not ideal
 - Closing the connection may impact other users
- Server.R
 - One connection per user session
 - Slower initial application load
 - No interference with other users

Make sure to close the connection when we are finished!

Storing DB connection details

- Don't want to store sensitive information directly in code!
- Use environment variables instead
- Stored in a .env file in format:
VAR_NAME = value
- In global.R:
 - `readRenviron(".env")`
- Using environment variables (watch out for object types):
 - `Sys.getenv("VAR_NAME")`

Query types

- **SELECT [cols] from [table]**
 - Retrieving rows from a specified database table where some conditions are met
 - ***Does not make any changes to the data***
- INSERT INTO [table]([cols])
 - Add new rows to a specified database table, with the provided column values
- UPDATE [table]
 - Modify existing rows in a specified database table
- DELETE FROM [table]
 - Remove rows from a specified database table where some conditions are met

Integrating to our app: Sample query

```
getTableName <- function() {  
  
  query <- 'SELECT  
column1, [column2]  
FROM schema.tableName  
WHERE column1 = `var` '  
  result <- dbGetQuery(dbPool, query)  
}
```

[PostgreSQL Query Cheat Sheet](#)

Integrating to our app: server.R

```
tableData <- reactive({  
  getTableName()  
})
```

```
filteredData <- reactive({  
  req(tableData())  
  ...  
})
```

```
onStop(function() { dbDisconnect(dbConn) })
```

Parameterised SQL queries

- SQL queries often need to be customised to accept different parameters
- In Shiny, we often need SQL queries based on user inputs
- Parameterised SQL queries let us make query templates, then fill in with variables when we evaluate the query
- Need to do this carefully to avoid SQL injection!

```
# Not parameterised
query <- "SELECT * FROM world_temp_data
        WHERE iso_a2 IN ('CA');"
```

```
#the parameter in question|
var <- "CA"
```

```
# dangerously parameterised
query <- paste0(
  "SELECT * FROM world_temp_data
  WHERE iso_a2 IN ('", var, "');"
)
```

```
# safely parameterised
query <- sqlInterpolate(
  "SELECT * FROM world_temp_data
  WHERE iso_a2 IN ?tag",
  tag = var
)
```

Safely parameterised SQL queries

- To protect against SQL injection, we can use the *sqlInterpolate* function
- This allows us to include variables safely in our SQL queries
- Protects against malicious user inputs
- Use ***dbExecute(con, query)*** or ***dbGetQuery(con, query)*** to process

```
sql <- "SELECT * FROM X WHERE name = ?name"
sqlInterpolate(ANSI(), sql, name = "Nick")

# This is safe because the single quote has been double escaped
sqlInterpolate(ANSI(), sql, name = "H'); DROP TABLE--;")

# Using paste0() could lead to dangerous SQL with carefully crafted inputs
# (SQL injection)
name <- "H'); DROP TABLE--;"
paste0("SELECT * FROM X WHERE name = '", name, "'")

# Use SQL() or dbQuoteIdentifier() to avoid escaping
sql2 <- "SELECT * FROM ?table WHERE name in ?names"
sqlInterpolate(ANSI(), sql2,
               table = dbQuoteIdentifier(ANSI(), "X"),
               names = SQL("('a', 'b')"))
)

# Don't use SQL() to escape identifiers to avoid SQL injection
sqlInterpolate(ANSI(), sql2,
               table = SQL("X; DELETE FROM X; SELECT * FROM X"),
               names = SQL("('a', 'b')"))
)
```


dbplyr - dplyr for SQL queries

- Using dbplyr, we can read data directly from our SQL tables
- Automatically loaded when dplyr detects database connectivity
- Use regular dplyr syntax to transform data
- ‘Lazy’ evaluation – SQL code is only evaluated in the database when the data is needed (like reactivity)
- Some more complex data transformations may be hard to replicate in dbplyr



dbplyr – dplyr for SQL queries

- Create a reference to the database table using ***tbl(con, table)***
- Write dplyr data transformations as usual
- When this code runs, dbplyr generates a SQL query (visible with **`show_query()`**) and evaluates against the database table
- Use the **`collect()`** function to retrieve results

```
> query <- tbl(con, "world_temp_data") %>%  
+   filter(iso_a2 %in% "CA") %>%  
+   select(country, city, region, month, year)  
> show_query(query)  
<SQL>  
SELECT "country", "city", "region", "month", "year"  
FROM "world_temp_data"  
WHERE ("iso_a2" IN ('CA'))  
> collect(query)  
# A tibble: 74,245 × 5  
  country city      region      month  year  
  <chr>   <chr>   <chr>      <int> <int>  
1 Canada  Calgary North America    1  1995  
2 Canada  Calgary North America    1  1995  
3 Canada  Calgary North America    1  1995  
4 Canada  Calgary North America    1  1995  
5 Canada  Calgary North America    1  1995
```

Performance considerations

Improving performance

Think of a Shiny app like a restaurant:

- The R process is represented by the chef
- User requests are customer's meal orders
- Different approaches to optimising the restaurant:
 - **Make the chef more efficient (optimise R code)**
 - Hire a prep cook (pre-processing)
 - Hire more chefs (add more R processes)
 - Purchase more equipment (more server memory)
 - Construct a new restaurant (multiple servers)

How to fix our pain points

- Minimise access to files / databases
- Calculations should be done once
- Don't Repeat Yourself (DRY principle)
- Filter datasets by most common feature first
- Separate filtering into staged reactivities

What pain points have you run into?

Performance tips

- Preprocessing
- Profiling
- Caching
- Computer resources
- Network monitoring (in browser)
- Load testing

Reactive functions and caching

- Reactive functions in Shiny already enable some caching
- When a reactive is called first time, results are stored in memory
- Next time that reactive is called, results are looked up
- Change in any reactive values in expression recalculates the reactive
- *What if we want to cache different combinations of those values?*

bindCache

- Allows reactive outputs to be cached in memory for quick access
- Reactive variables form the 'key' for that reactive
- Different levels of caching available:
 - Session – one cache per user session.
 - Application – cache is shared across multiple sessions. Users can take access calculations done in another user's session

```
city_data <- reactive({  
  fetchData(input$city)  
}) %>%  
bindCache(input$city)
```

Exercise

In RStudio Cloud, open **Session-5**, then `/stage1`

- Create a `.env` file (using `.env.example` as a template), fill it with DB connection details
- In `global.R`
 - Load the `.env` file in `global.R` using `readRenvron()`, then use the environment variables with `Sys.getenv()` to create a db connection
 - Complete the `getCountries` function, using `dbplyr` to query the `world_temp_data` table from the database, filter by country `iso_a2` codes, and summarise an average monthly temperature
- In `temperature-page-module.R`,
 - Create a reactive to store the `iso_a2` codes from `filtered_data`
 - Create a reactive to retrieve temperature data from DB based on the selected countries using your `getCountries` function
 - Create a reactive to filter the country temperature data by year, then use this reactive in the `temperature_table`
 - Apply `bindCache` as appropriate to these reactives to improve performance.

Temperature data structure

	id	region	country	iso_a2	city	month	day	year	date	avg_temp
1	1089209	North America	Canada	CA	Calgary	1	1	1995	1995-01-01	12.6
2	1089210	North America	Canada	CA	Calgary	1	2	1995	1995-01-02	4.5
3	1089211	North America	Canada	CA	Calgary	1	3	1995	1995-01-03	2.5
4	1089212	North America	Canada	CA	Calgary	1	4	1995	1995-01-04	11.4
5	1089213	North America	Canada	CA	Calgary	1	5	1995	1995-01-05	11.3
6	1089214	North America	Canada	CA	Calgary	1	6	1995	1995-01-06	4.0
7	1089215	North America	Canada	CA	Calgary	1	7	1995	1995-01-07	5.4
8	1089216	North America	Canada	CA	Calgary	1	8	1995	1995-01-08	4.5
9	1089217	North America	Canada	CA	Calgary	1	9	1995	1995-01-09	7.1
10	1089218	North America	Canada	CA	Calgary	1	10	1995	1995-01-10	17.1

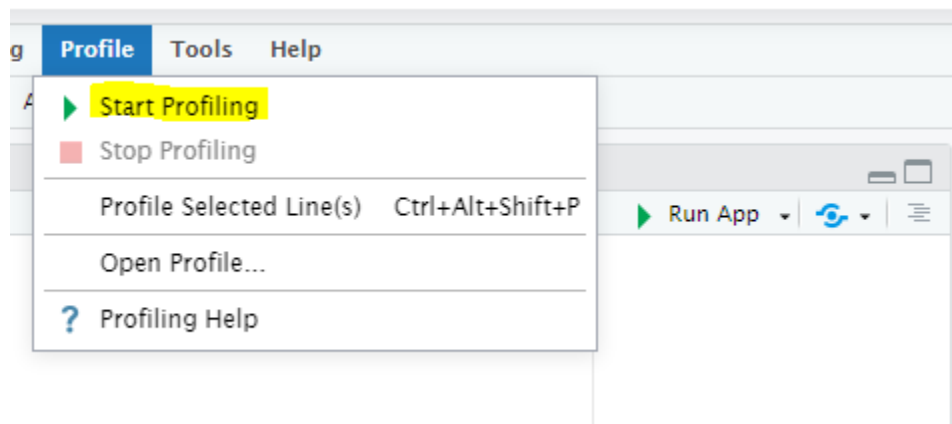
Profiling

Profiling / Profvis

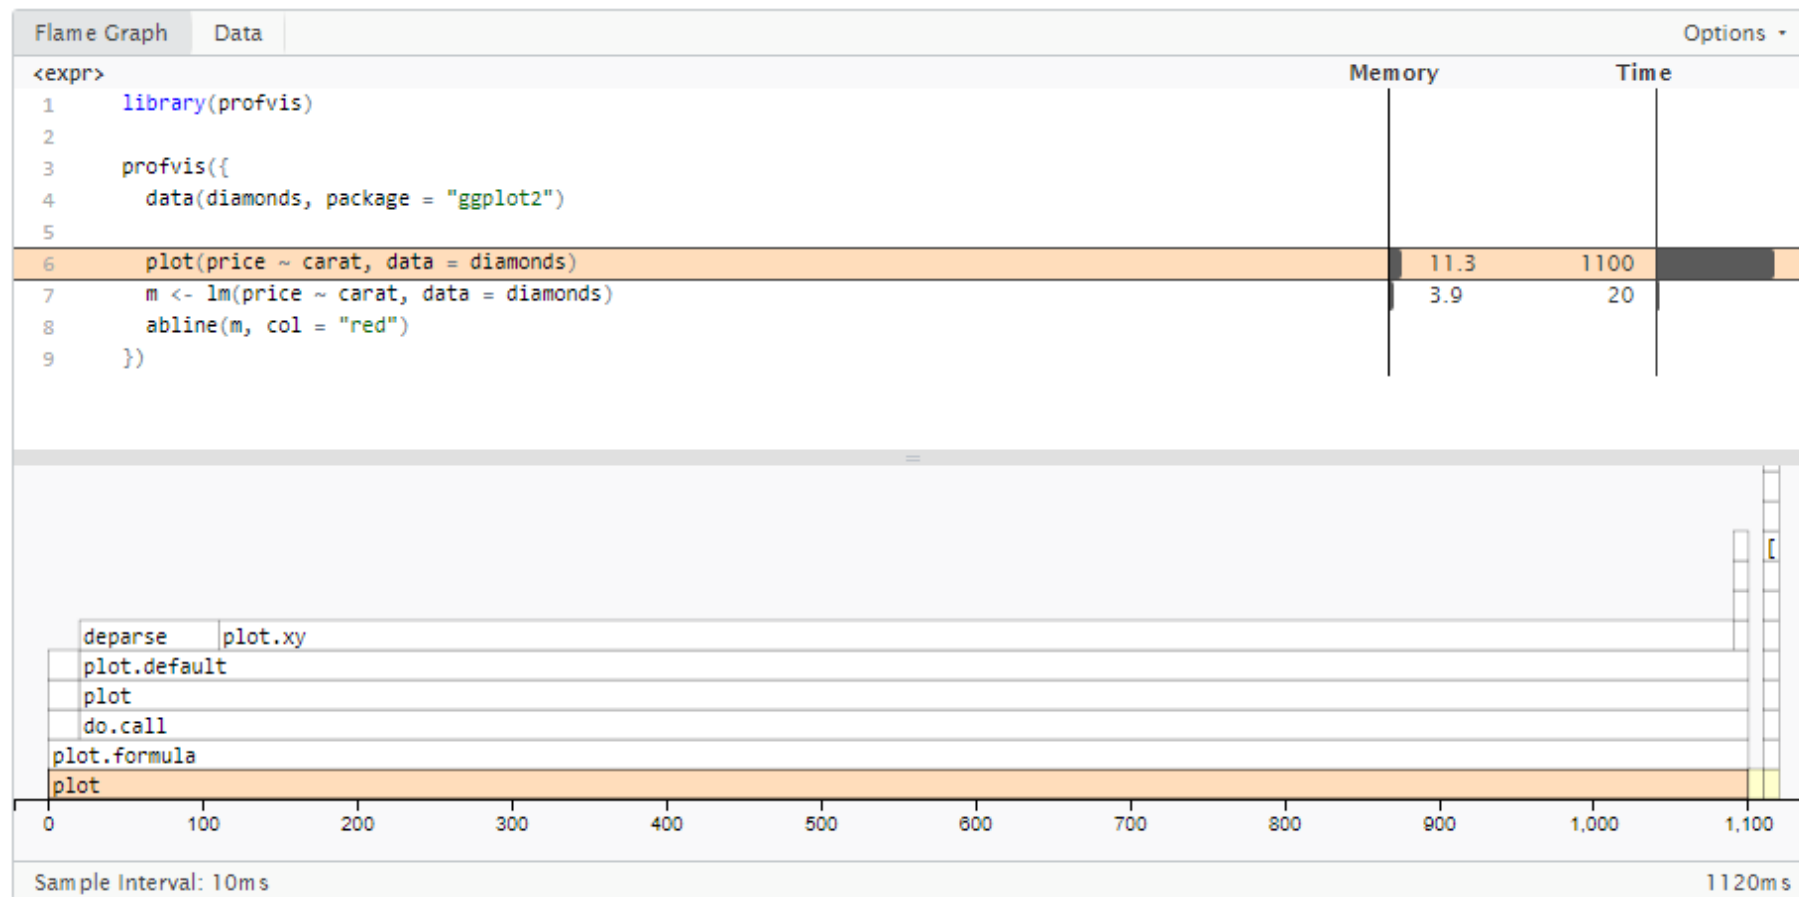
- Built into RStudio IDE
- Records app performance over time
- Allows you to find your pain points
- NOTE: Only optimize if needed

Profiling: How to

- Profile > Start Profiling
 - This can be done before or during the Shiny app runtime
- Take actions in the Shiny app to make the server do some work
- Profile > Stop Profiling



Results: flame graph



Next time

- Profiling
- Automated Reports with Shiny and Quarto

Challenge:

- Move your connection into *server.R* so you only make one connection per user
- Rework the *getCountries()* method to include the year as part of the query
- Change the *getCountries()* method to write a *sqlInterpolate* query instead of using *dbplyr*
- Profile the app to identify other performance areas to focus on, see how you can improve the performance further
- Share your work on the session 5 forum!